

vir-projekt

July 31, 2016

Contents

1	Convolutional Network Features	1
1.1	Setup der Programmierungsumgebung	1
1.1.1	Python-Libraries	1
1.1.2	mAP	2
1.2	Vorverarbeitung der Daten	3
1.2.1	Trainingsdaten erzeugen	3
1.2.2	Trainings und Evaluierungsset erstellen	4
2	VGG16 mit Keras laden	5
3	Visualisierung der Filter des VGG Models	7
4	Feature-Vektoren	9
4.1	Modelkonstruktion 1	10
4.2	Trainingsergebniss	10
4.3	Genauigkeit	11
4.4	Modelkonstruktion 2	12
5	Evaluierung	13
5.1	Visualisierung der Erkennungsleistung	14
5.2	Precision-Recall	17

1 Convolutional Network Features

1.1 Setup der Programmierungsumgebung

1.1.1 Python-Libraries

```
In [1]: %matplotlib inline
import scipy.misc as sm
import numpy as nd
import glob as g
import os.path as op
import re
import scipy.spatial.distance as sdist
import numpy as np
```

```

import matplotlib.pyplot as plt
from skimage import io, transform
import pandas as pd

import pickle

import seaborn as sns
import functools
import itertools
from collections import namedtuple
import random
from sklearn.manifold import TSNE
from sklearn.decomposition import PCA
import h5py
import random
from tqdm import tqdm

```

```

In [2]: from keras.layers import Input, Dense, Lambda, Dropout, Flatten, MaxPooling2D
        from keras.layers import Convolution2D, MaxPooling2D, ZeroPadding2D
        from keras.models import Model, Sequential
        from keras import backend as K
        from keras import objectives
        from keras.datasets import mnist
        from keras.optimizers import SGD, Adam

```

Using Theano backend.

1.1.2 mAP

```

In [3]: def feature_distances(query, features, f_distance=sdist.cityblock):
        return np.apply_along_axis(f_distance, 1, features, query)

def p_r_curve(X, Y, f_distance=sdist.cityblock, query_dist=feature_distances):
    pres, recal = np.zeros(len(X)), np.zeros(len(X))
    ap = np.zeros(len(X))
    for q in range(len(X)):
        distances = query_dist(X[q], X, f_distance)
        retrieval = (Y == Y[q])[distances.argsort()]
        p_q = retrieval.cumsum() / (np.arange(len(retrieval)) + 1)
        pres = np.add(pres, p_q)
        recal = np.add(recal, retrieval.cumsum() / retrieval.sum())
        ap[q] = p_q[retrieval].mean()
    return recal / len(X), pres / len(X), ap

In [66]: img_shape = 3, 224, 224
        nb_training_images = 2
        angle_range = 45

```

```
crop_range = 20

sns.set_style("whitegrid", {'axes.grid' : False})
sns.set_context('talk')
```

1.2 Vorverarbeitung der Daten

- BildgröSse skalieren, da NN-Eingangsvektor konstant
- Anordnen: RGBRGBRGB -> RRR...GGG...BBB

```
In [5]: def reorder(img):
        img = img * 255
        img[:, :, 0] -= 103.939
        img[:, :, 1] -= 116.779
        img[:, :, 2] -= 123.68
        img = img.transpose((2, 0, 1))
        return img

def preprocess(img, img_shape=(224, 224, 3)):
    return reorder(transform.resize(img, img_shape))
```

```
In [6]: paths = sorted(g.glob('./data/WebImages_71x6/*.jpg'))
        data = pd.DataFrame([(re.split('_', op.basename(path))[0], path)
                             for path in paths], columns=['class', 'path'])
        classes = list(enumerate(sorted(list(set(data['class'])))))
        class_map = {name: index for index, name in classes}
        data['y'] = [class_map[name] for name in data['class']]
        data.head()
```

```
Out[6]:   class      path  y
0     50  ./data/WebImages_71x6/50_1.jpg  0
1     50  ./data/WebImages_71x6/50_2.jpg  0
2     50  ./data/WebImages_71x6/50_3.jpg  0
3     50  ./data/WebImages_71x6/50_4.jpg  0
4     50  ./data/WebImages_71x6/50_5.jpg  0
```

1.2.1 Trainingsdaten erzeugen

Die Trainingsdaten pro Klasse bestehen aus zwei Bildern aus dem Datenset. Aus jedem Bild werden 10 weitere Bilder erzeugt.

```
In [7]: import warnings
        warnings.filterwarnings('ignore')


In [38]: def rand_img(img):
        height, width, depth = img.shape
        if random.choice((True, False)):
            img = np.fliplr(img)
```

```

img = img[random.randint(0,crop_range):
           random.randint(height - crop_range,height),
           random.randint(0,crop_range):
           random.randint(width - crop_range,width)]
img = transform.rotate(img,(random.random()-0.5) * 2 * angle_range,
                       cval=127, resize=False,mode='constant',order=2)

return img
plt.imshow(rand_img(io.imread(data['path'].ix[200])), interpolation='none')
plt.show()

```



1.2.2 Trainings und Evaluierungsset erstellen

- Trennung von Trainingsdaten und Testdaten
- Arbeit mit NN macht sonst keinen Sinn

```

In [39]: x_train_df = data[:0]
         x_test_df  = data[:0]
         for key, group in data.groupby('class'):
             x_train_df = x_train_df.append(group[0:nb_training_images])
             x_test_df  = x_test_df.append(group[nb_training_images:])
         x_train_df.head()

def build_np(df, augment=0):

```

```

nb_images = len(df) + len(df) * augment
x = np.zeros((nb_images,) + img_shape, dtype=np.float32)
y = np.zeros((nb_images,), dtype=np.int8)
idx = 0
# print(x.shape)
for path, cId in tqdm(list(zip(df['path'], df['y']))):
    img = io.imread(path)
    x[idx] = preprocess(img)
    y[idx] = cId
    idx += 1
    for i in range(augment):
        aug_img = rand_img(img)
        x[idx] = preprocess(aug_img)
        y[idx] = cId
        idx += 1
numpy = np.zeros((len(y), max(y) + 1), dtype=np.bool)
numpy[np.arange(len(y)), y] = 1
return x, numpy

x_train, y_train = build_np(x_train_df, augment=10)
x_test, y_test = build_np(x_test_df)

```

```

100%| 142/142 [00:37<00:00, 3.91it/s]
100%| 284/284 [00:02<00:00, 131.28it/s]

```

```

In [ ]: with open('data/dataset.pickle', 'rb') as feat:
        x_train, y_train, x_test, y_test = pickle.load(feat)

```

2 VGG16 mit Keras laden

siehe <https://arxiv.org/abs/1409.1556>:

Very Deep Convolutional Networks for Large-Scale Image Recognition
 Karen Simonyan, Andrew Zisserman

Parameter und Quellcode basierend auf diesem [Gist](#)

Abwandlung hier: nur eine bestimmte Anzahl von Layern werden geladen. Die Ausgabe des Netzes an diesem Punkt, kann dann als Input für eine neues Netz genutzt werden, wie es in [CS231n](#) beschrieben wurde.

```

In [13]: def load_some_weights(model, weights_path):
        f = h5py.File(weights_path, mode='r')
        nb_layers = f.attrs['nb_layers']
        for k in range(len(model.flattened_layers)):

```

```

        g = f['layer_{}'.format(k)]
        weights = [g['param_{}'.format(p)] for p in range(g.attrs['nb_params'])]
        model.flattened_layers[k].set_weights(weights)

    f.close()
    return model

def VGG_16(weights_path=None, layer_limit=None):
    model = Sequential()
    model.add(ZeroPadding2D((1,1),input_shape=(3,224,224)))
    model.add(Convolution2D(64, 3, 3, activation='relu'))
    model.add(ZeroPadding2D((1,1)))
    model.add(Convolution2D(64, 3, 3, activation='relu'))
    model.add(MaxPooling2D((2,2), strides=(2,2)))
    # img dim = 112
    model.add(ZeroPadding2D((1,1)))
    model.add(Convolution2D(128, 3, 3, activation='relu'))
    model.add(ZeroPadding2D((1,1)))
    model.add(Convolution2D(128, 3, 3, activation='relu'))
    model.add(MaxPooling2D((2,2), strides=(2,2)))
    # img dim = 56
    model.add(ZeroPadding2D((1,1)))
    model.add(Convolution2D(256, 3, 3, activation='relu'))
    model.add(ZeroPadding2D((1,1)))
    model.add(Convolution2D(256, 3, 3, activation='relu'))
    model.add(ZeroPadding2D((1,1)))
    model.add(Convolution2D(256, 3, 3, activation='relu'))
    model.add(MaxPooling2D((2,2), strides=(2,2)))
    # img dim = 28
    model.add(ZeroPadding2D((1,1)))
    model.add(Convolution2D(512, 3, 3, activation='relu'))
    model.add(ZeroPadding2D((1,1)))
    model.add(Convolution2D(512, 3, 3, activation='relu'))
    model.add(ZeroPadding2D((1,1)))
    model.add(Convolution2D(512, 3, 3, activation='relu'))
    model.add(MaxPooling2D((2,2), strides=(2,2)))
    # img dim = 14
    model.add(ZeroPadding2D((1,1)))
    model.add(Convolution2D(512, 3, 3, activation='relu'))
    model.add(ZeroPadding2D((1,1)))
    model.add(Convolution2D(512, 3, 3, activation='relu'))
    model.add(ZeroPadding2D((1,1)))
    model.add(Convolution2D(512, 3, 3, activation='relu'))
    model.add(MaxPooling2D((2,2), strides=(2,2)))
    # img dim = 7
    model.add(Flatten())
    # length 25088
    if layer_limit == 'maxpool4':

```

```

        return load_some_weights(model, weights_path)

    model.add(Dense(4096, activation='relu'))
    model.add(Dropout(0.5))
    model.add(Dense(4096, activation='relu'))
    model.add(Dropout(0.5))

    return load_some_weights(model, weights_path)

In [40]: f_model = VGG_16('models/vgg16/vgg16_weights.h5')
        sgd = SGD(lr=0.1, decay=1e-6, momentum=0.9, nesterov=True)
        f_model.compile(optimizer=sgd, loss='categorical_crossentropy')

In [41]: conv_model = VGG_16('models/vgg16/vgg16_weights.h5', layer_limit='maxpool4')
        sgd = SGD(lr=0.1, decay=1e-6, momentum=0.9, nesterov=True)
        conv_model.compile(optimizer=sgd, loss='categorical_crossentropy')

```


3 Visualisierung der Filter des VGG Models

Das VGG16 Netz benutzt nur 3x3-Filterkernel. Die erste Abbildung zeigt die 64 Filter des ersten Layers und die zweite Abbildung zeigt die 256 Filter des siebten Layers als Beispiele.


```

In [16]: conv1 = f_model.layers[1]
        img = np.zeros((8*3, 8*3))
        for y in range(8):
            for x in range(8):
                # print(y,x, y*8+x)
                img[x*3:x*3+3,y*3:y*3+3] = conv1.get_weights()[0][y*8 + x, :, :, 0]
        plt.imshow(img, interpolation='none')
        plt.show()

```



```
In [17]: conv2 = f_model.layers[15]
dim = 16
img = np.zeros((dim*3, dim*3))
for y_f in range(dim):
    for x_f in range(dim):
        for y in range(dim):
            for x in range(dim):
                # print(y,x, y*8+x)
                img[x*3:x*3+3,y*3:y*3+3] = conv2.get_weights()[0][y*dim + x,0,:,:]
plt.imshow(img, interpolation='none')
plt.show()
```

4 Feature-Vektoren

Für alle Trainings und Testdaten werden die Featurevektoren auf der Festplatte zwischengespeichert.

```
In [43]: dense_features_train = f_model.predict(x_train,verbose=1)
dense_features_test  = f_model.predict(x_test,verbose=1)
with open('models/vgg_dense.features', 'wb') as feat:
    pickle.dump((dense_features_train, dense_features_test), feat)

conv_features_train = conv_model.predict(x_train,verbose=1)
conv_features_test  = conv_model.predict(x_test,verbose=1)
with open('models/vgg_conv.features', 'wb') as feat:
    pickle.dump((conv_features_train, conv_features_test), feat)

1562/1562 [=====] - 920s
284/284 [=====] - 194s
1562/1562 [=====] - 1029s
```

284/284 [=====] - 174s

```
In [18]: with open('models/vgg_dense.features', 'rb') as feat:
         dense_features_train, dense_features_test = pickle.load(feat)
         with open('models/vgg_conv.features', 'rb') as feat:
             conv_features_train, conv_features_test = pickle.load(feat)
```

4.1 Modelkonstruktion 1

Zuerst wurde versucht zwei Netzwerke auf Basis des CNN-Tensors(7, 7, 512) als Eingabe zu trainieren.

```
In [81]: t_model = Sequential()
         t_model.add(Dense(71, input_shape=(25088,), activation='softmax'))
         adam1 = Adam(lr=0.00003)
         t_model.compile(optimizer=adam1,
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])

         t_model_multi = Sequential()
         t_model_multi.add(Dense(4096, input_shape=(25088,), activation='relu'))
         t_model_multi.add(Dropout(0.5))
         t_model_multi.add(Dense(4096, activation='relu'))
         t_model_multi.add(Dropout(0.5))
         t_model_multi.add(Dense(71, activation='softmax'))
         adam2 = Adam(lr=0.00003)
         t_model_multi.compile(optimizer=adam2,
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])
```

4.2 Trainingsergebniss

Sehr zeitaufwendig ohne GPU. Nur das Netzwerk mit 3 Layern war erfolgsversprechend und wurde länger trainiert.

```
In [90]: def train_model(model, n=1):
         history = model.fit(conv_features_train, y_train,
                        verbose=0,
                        shuffle=True,
                        nb_epoch=10,
                        batch_size=32)

         loss, acc = model.evaluate(conv_features_test, y_test, verbose=0)
         return history, loss, acc


In [82]: t_model_eval = train_model(t_model)
```

```
In [91]: batched_training = []
        for idx in tqdm(range(12)):
            batched_training.append(train_model(t_model_multi,n=1))
```

100%|| 12/12 [4:51:01<00:00, 2135.03s/it]

```
In [105]: line, = plt.plot(t_model_eval[0].history['loss'], label="1 FC")
        line, = plt.plot(h, label="3 FC")
        plt.legend()
```

Out[105]: <matplotlib.legend.Legend at 0x7f8c43bc4c50>



4.3 Genauigkeit

```
In [111]: evaluations = []
        loss, acc = t_model.evaluate(conv_features_test,y_test, verbose=1)
        evaluations.append(('1 FC', loss, acc))
        loss, acc = t_model_multi.evaluate(conv_features_test,y_test, verbose=1)
        evaluations.append(('3 FC', loss, acc))
        pd.DataFrame(evaluations, columns=['Name', 'Loss', 'Genauigkeit'])
```

284/284 [=====] - 0s

284/284 [=====] - 0s

```
Out[111]:
```

	Name	Loss	Genauigkeit
0	1 FC	12.929887	0.084507
1	3 FC	6.155772	0.556338


4.4 Modelkonstruktion 2

Dieses Model verwendet die Ausgabe des Fully Connected Layers als Feature (4096 Werte)

```
In [47]: dense_model = Sequential()
dense_model.add(Dense(71, input_shape=(4096,), activation='softmax'))
adam = Adam(lr=0.000006)
dense_model.compile(optimizer=adam,
                    loss='categorical_crossentropy',
                    metrics=['accuracy'])
history = dense_model.fit(dense_features_train, y_train,
                          verbose=0,
                          shuffle=True,
                          nb_epoch=150,
                          batch_size=16)

loss, acc = dense_model.evaluate(dense_features_test, y_test, verbose=0)
print(acc)
plt.plot(history.history['loss'])
plt.show()
```


0.823943660293



5 Evaluierung


Zur Evaluierung wurde zuerst eine [PCA](#) durchgeführt. Danach wurde diese reduzierte Parameterraum mittels [t-SNE](#) visualisiert. Die erste Abbildung zeigt die Trainingsdaten im neuen Parameterraum und die Zweite zeigt die Testdaten. Die Cluster in der ersten Abbildung entstehen wahrscheinlich durch die erzeugten Trainingsdaten.

```
In [75]: preds = dense_model.predict(dense_features_train)
pca = PCA(n_components=50)
pca.fit(preds)
tsne = TSNE(n_components=2, random_state=0)
np.set_printoptions(suppress=True)
transformed = tsne.fit_transform(pca.transform(preds))
data = pd.DataFrame(transformed, columns=['x', 'y'])
data['class'] = pd.Series(x_test_df['y'])
g = sns.lmplot(x='x', y='y', data=data, hue='class', fit_reg=False, legend=False)
```



```
In [76]: preds = dense_model.predict(dense_features_test)
pca.fit(preds)
np.set_printoptions(suppress=True)
transformed = tsne.fit_transform(pca.transform(preds))
data = pd.DataFrame(transformed, columns=['x', 'y'])
data['class'] = pd.Series(x_test_df['y'])
sns.lmplot(x='x', y='y', data=data, hue='class', fit_reg=False, legend=False)
```

```
Out[76]: <seaborn.axisgrid.FacetGrid at 0x7f8c46190f60>
```



5.1 Visualisierung der Erkennungsleistung

Um die Erkennungsleistung des Netzwerks zu visualisieren wurde die [Sensibilität gegenüber Verdeckungen](#) gemessen.

```
In [50]: def gray_slide(test_img, y_test_img, n = 200, r = 20, step=5):
coordinates = []
for height, i in enumerate(range(r, test_img.shape[0] - r, step)):
    for width, j in enumerate(range(r, test_img.shape[1] - r, step)):
```

```

        coordinates.append((i,j))
    occluded = np.zeros((len(coordinates),) + img_shape, dtype=np.float32)
    idx = 0
    for idx in range(len(coordinates)):
        i, j = coordinates[idx]
        occ = test_img.copy()
        occ[i-r:i+r, j - r: j + r] = 127
        occluded[idx] = preprocess(occ)
    x_o_class = dense_model.predict(f_model.predict(occluded))
    z = x_o_class[:,y_test_img]
    return np.reshape(z - z.mean(),(height + 1, width + 1))

```

```
In [305]: gray_slide_img = gray_slide(io.imread(x_test_df['path'][n]), x_test_df['y'][n])
```

```
In [51]: slide_image_paths = random.sample(list(zip(x_test_df['path'], x_test_df['y'])), 5)
```

```
In [52]: slide_images = [gray_slide(io.imread(path), y) for path, y in tqdm(slide_image_paths)]
```

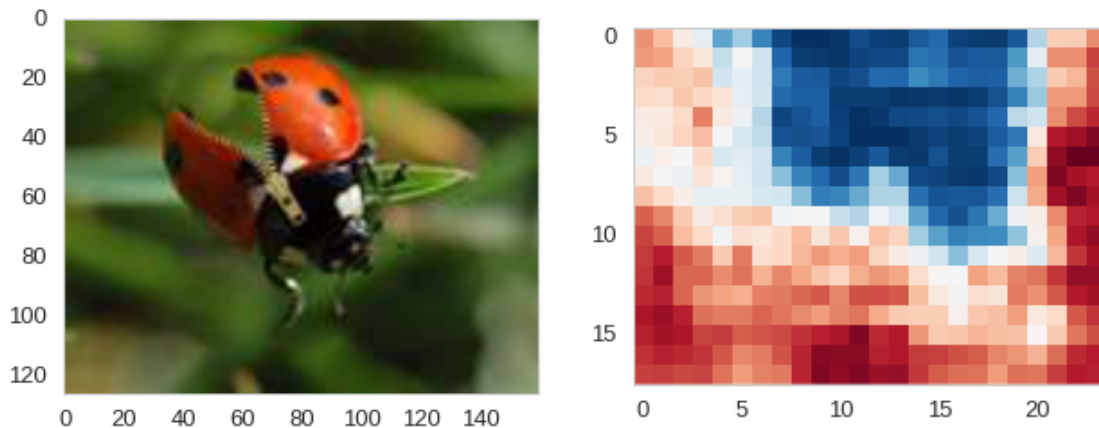
```
100%|| 5/5 [24:13<00:00, 300.25s/it]
```

```
In [336]: (slide_image_paths[0])
```


```
Out[336]: ('./data/WebImages_71x6/himbeere_6.jpg', 25)
```

```
In [77]: def plot_side_by_side(i):
    plt.subplot(121)
    plt.imshow(io.imread(slide_image_paths[i][0]))
    plt.subplot(122)
    plt.imshow(slide_images[i], cmap='RdBu_r', interpolation='none')
```


```
In [71]: plot_side_by_side(0)
```




```
In [70]: plot_side_by_side(1)
```



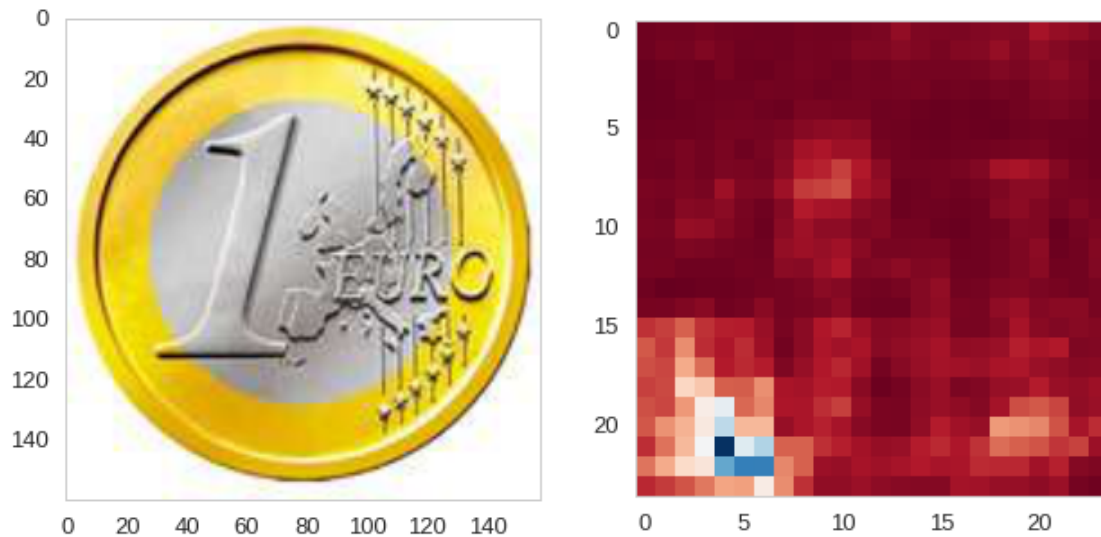
```
In [69]: plot_side_by_side(2)
```



```
In [68]: plot_side_by_side(3)
```




```
In [72]: plot_side_by_side(4)
```



5.2 Precision-Recall

```
In [112]: x_features = dense_model.predict(dense_features_test)
          recal, pres, mAP = p_r_curve(x_features,np.array(x_test_df['y']),sdist.euclidean)
          plt.plot(recal, pres,'.-g',label='FC feature {:.2f}'.format(mAP.mean()))

          x_features = t_model.predict(conv_features_test)
          recal, pres, mAP = p_r_curve(x_features,np.array(x_test_df['y']),sdist.euclidean)
          plt.plot(recal, pres,'.-b', label='CONV feature {:.2f}'.format(mAP.mean()))

          plt.xlim([0,1])
          plt.ylim([0,1])
          mAP.mean()
          plt.legend()
```

```
Out [112]: <matplotlib.legend.Legend at 0x7f8c43a4fa90>
```

